



Transformer Models

김윤식


+

o

•

Introduction

- Machine translation, Machine text generation, etc.: Seq2Seq
- Previously used RNN, LSTM, ...



- New approach: “Attention”

Overview

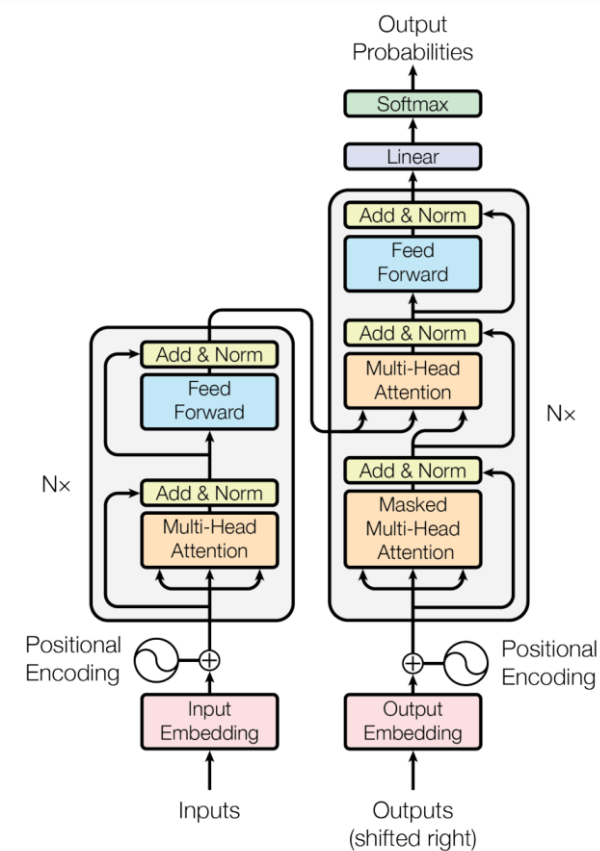
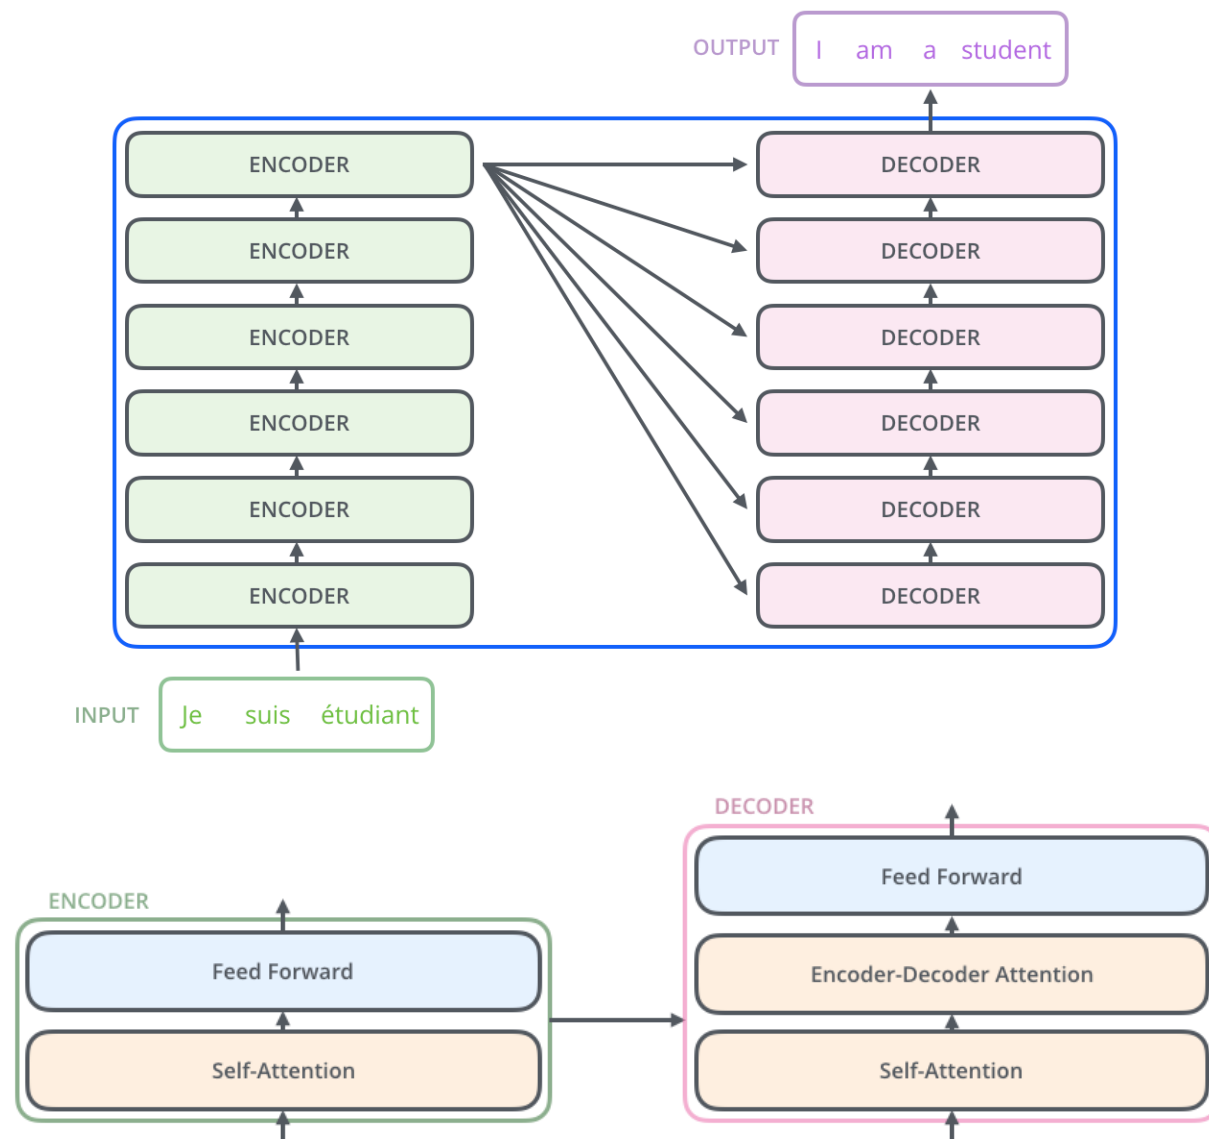
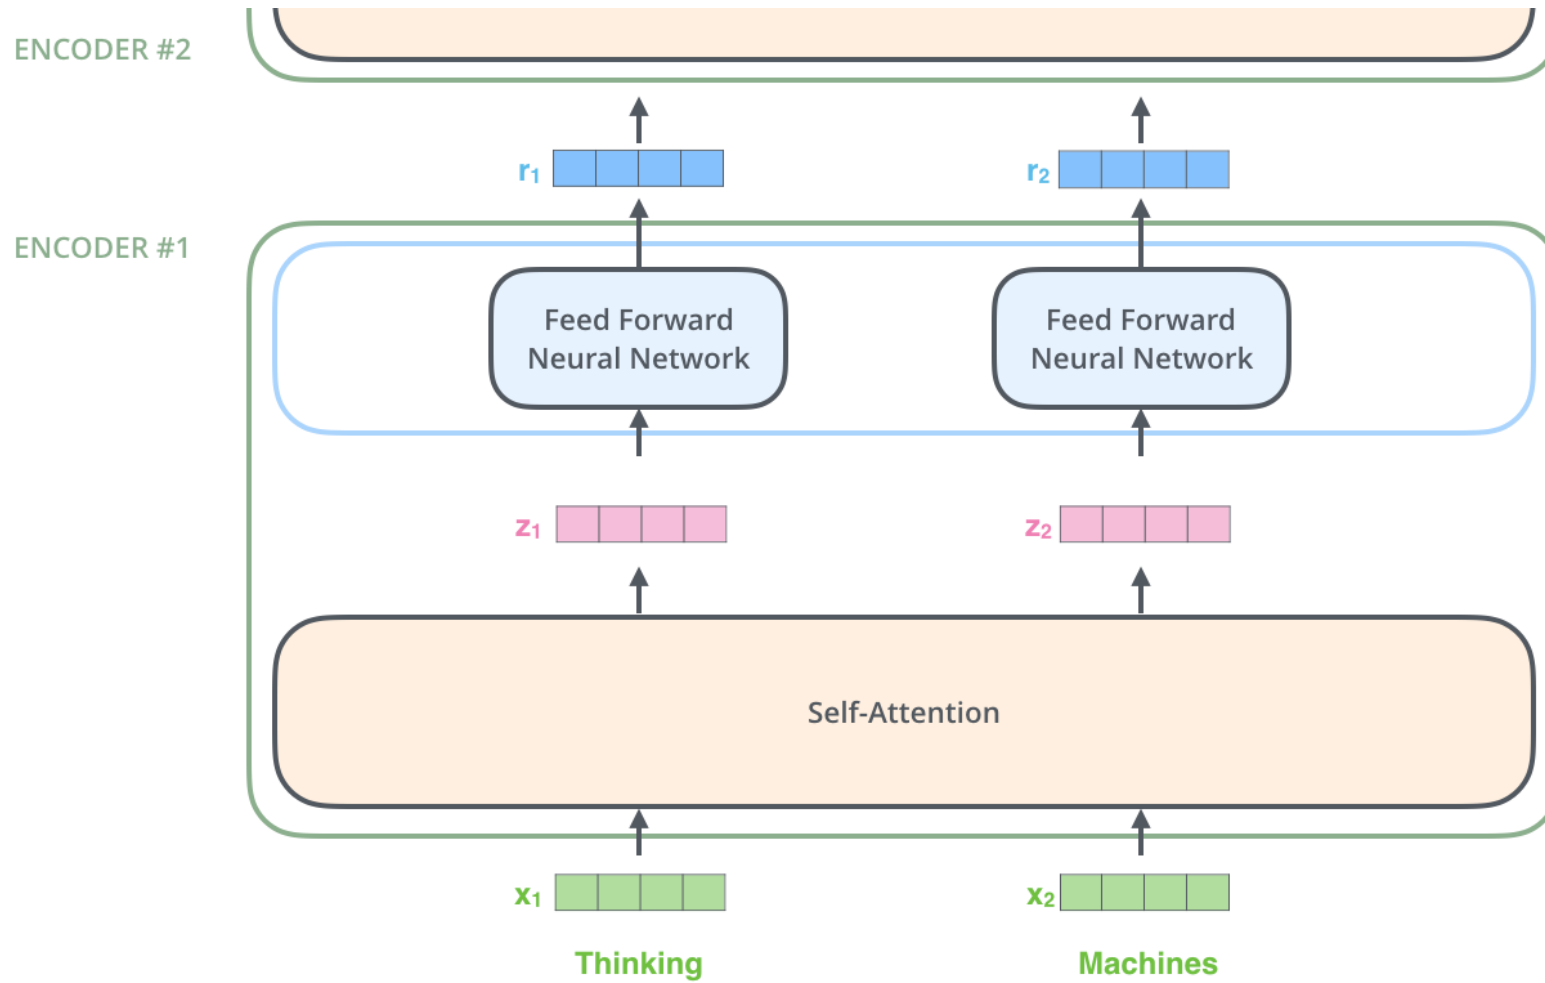
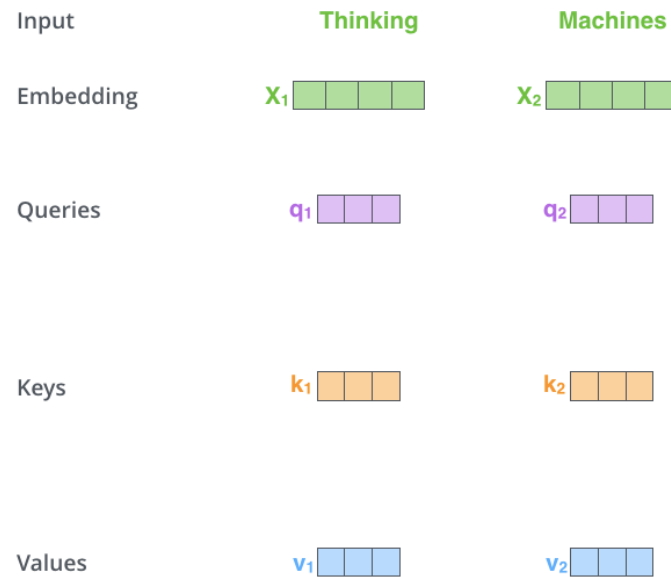


Figure 1: The Transformer - model architecture.

Encoder



Self-attention



Input

Embedding

Queries

Keys

Values

Score

Divide by 8 ($\sqrt{d_k}$)

Softmax

Softmax
X
Value

Sum

Thinking

x_1

q_1

k_1

v_1

$q_1 \cdot k_1 = 112$

14

0.88

v_1

z_1

Machines

x_2

q_2

k_2

v_2

$q_1 \cdot k_2 = 96$

12

0.12

v_2

z_2

Multi-headed Self-attention

1) This is our input sentence*


2) We embed each word*

3) Split into 8 heads. We multiply X or R with weight matrices





4) Calculate attention using the resulting $Q/K/V$ matrices

5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer

Thinking
Machines






* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one



...

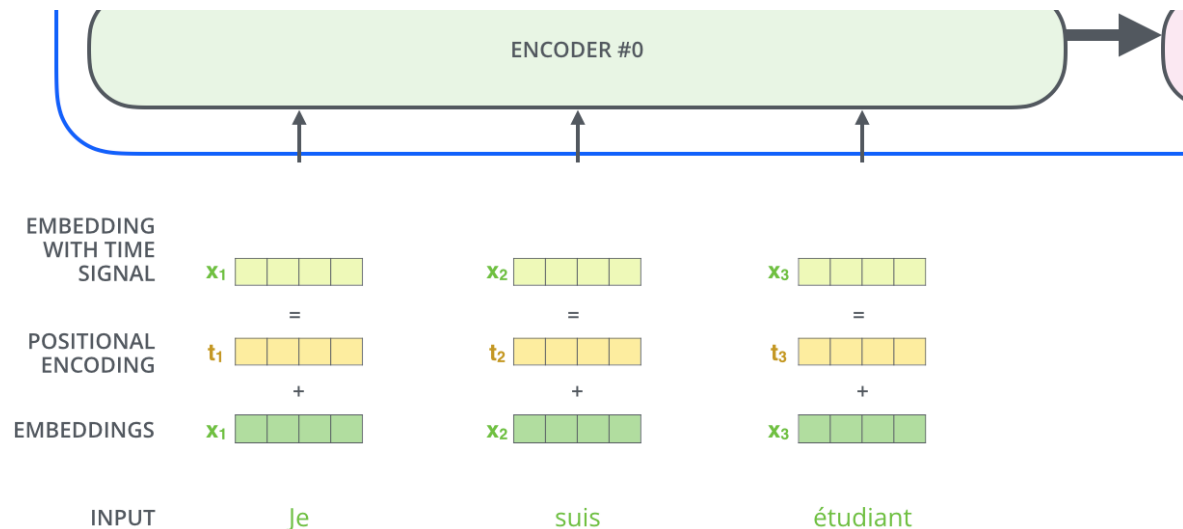
...

...




Positional Encoding

- Injecting information about the positions of tokens



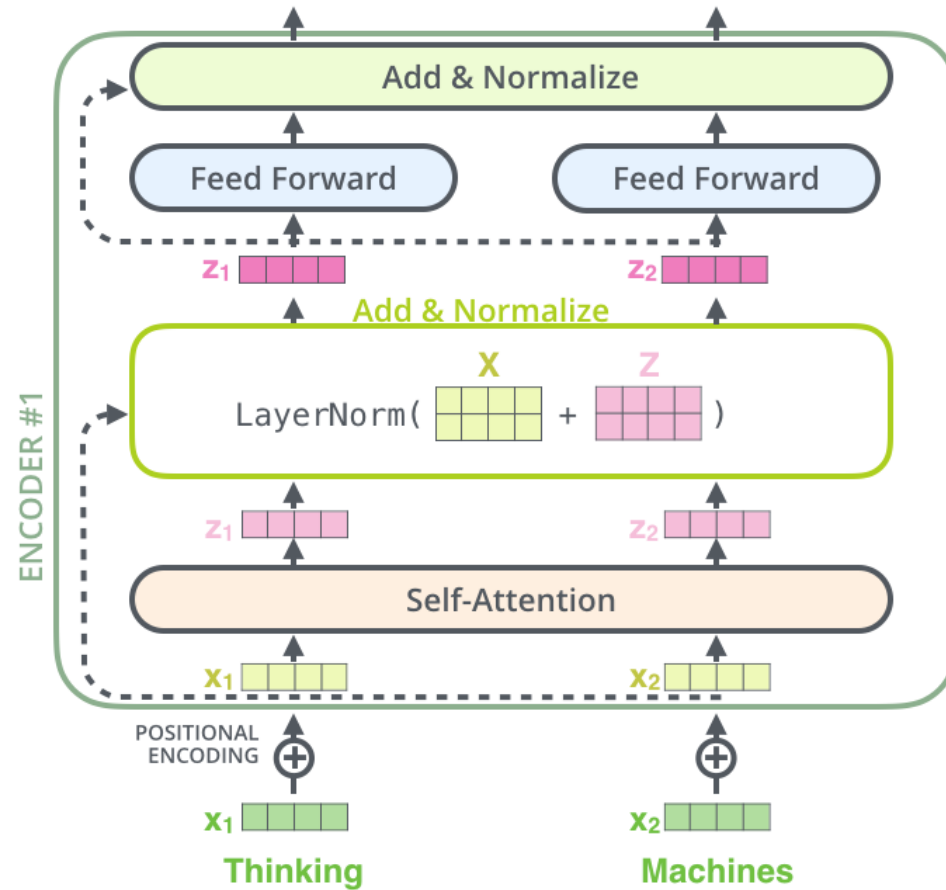
$$PE_{(pos,2i)} = \sin(pos/10000^{2i/d_{\text{model}}})$$

$$PE_{(pos,2i+1)} = \cos(pos/10000^{2i/d_{\text{model}}})$$

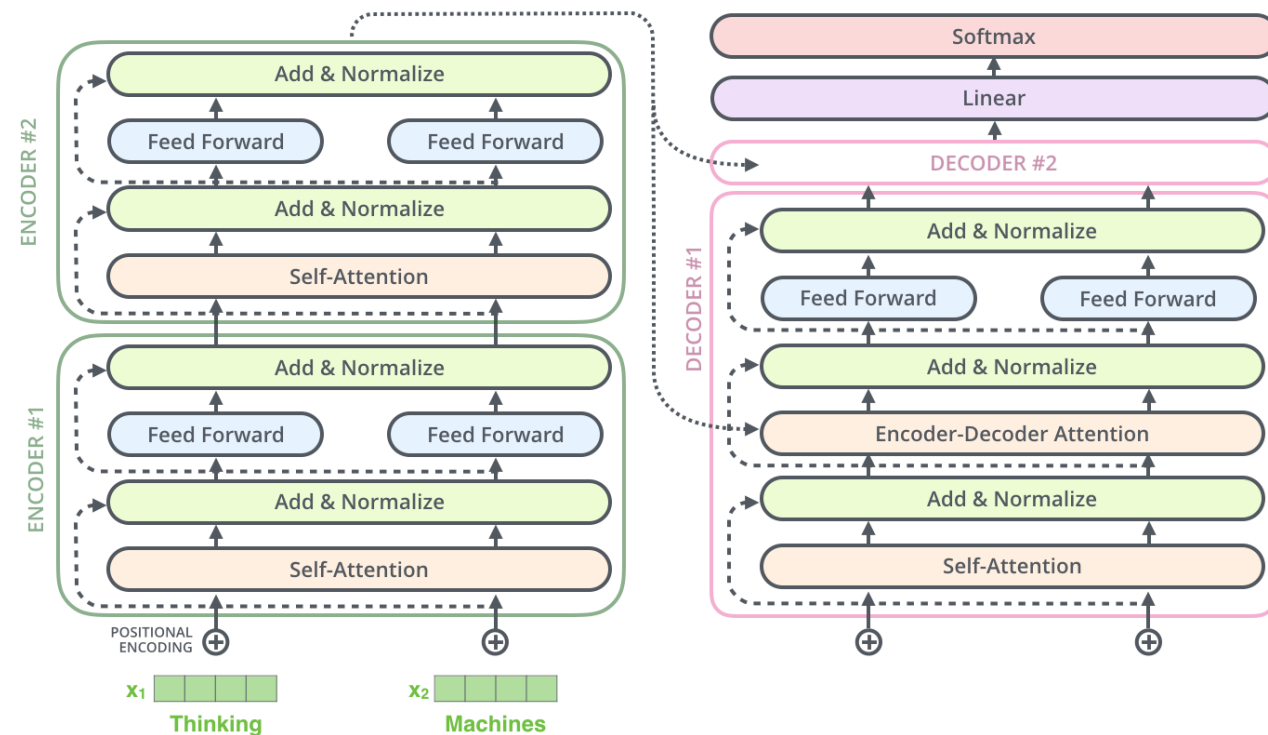
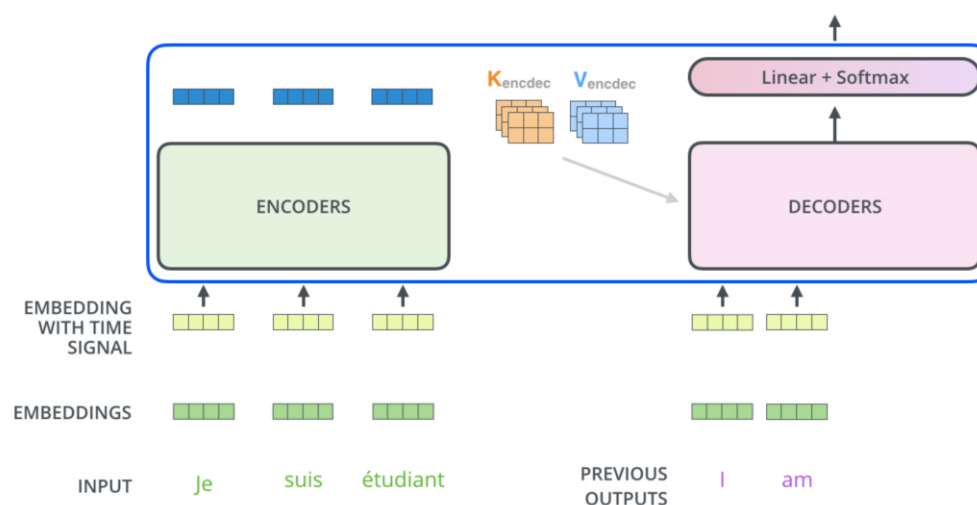
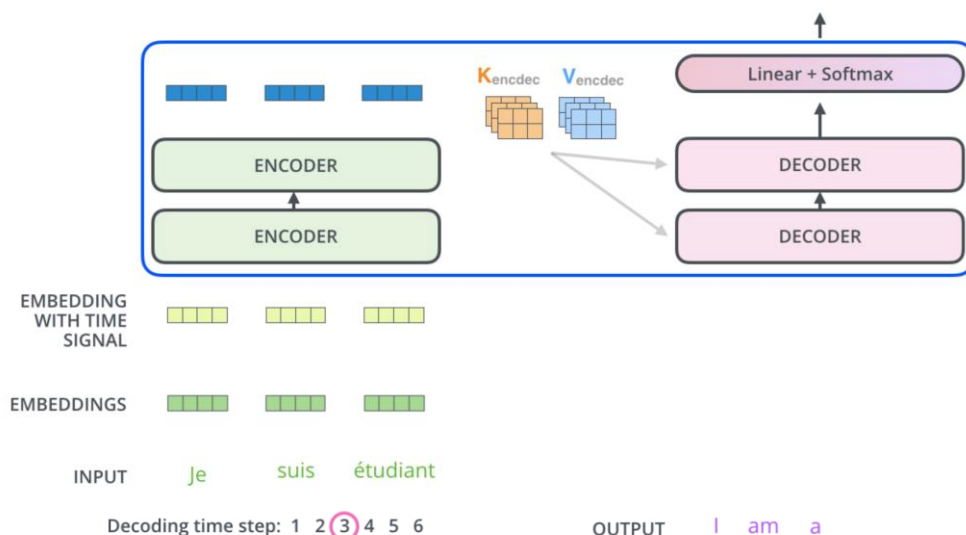


Overall structure of an Encoder

- Residual connections
- Layer normalization



Decoder



Encoder-Decoder Attention:

Same as Self-attention, but K and V vectors are from the encoder outputs


“Masked” Attention:

Since decoders should only consider earlier positions in the sequence, future positions are “masked” by setting them to $-\infty$ before the softmax step

Linear and Softmax layers

Which word in our vocabulary
is associated with this index?


Get the index of the cell
with the highest value
(**argmax**)



Training the model


Target Model Outputs

Output Vocabulary: a am I thanks student <eos>



Trained Model Outputs

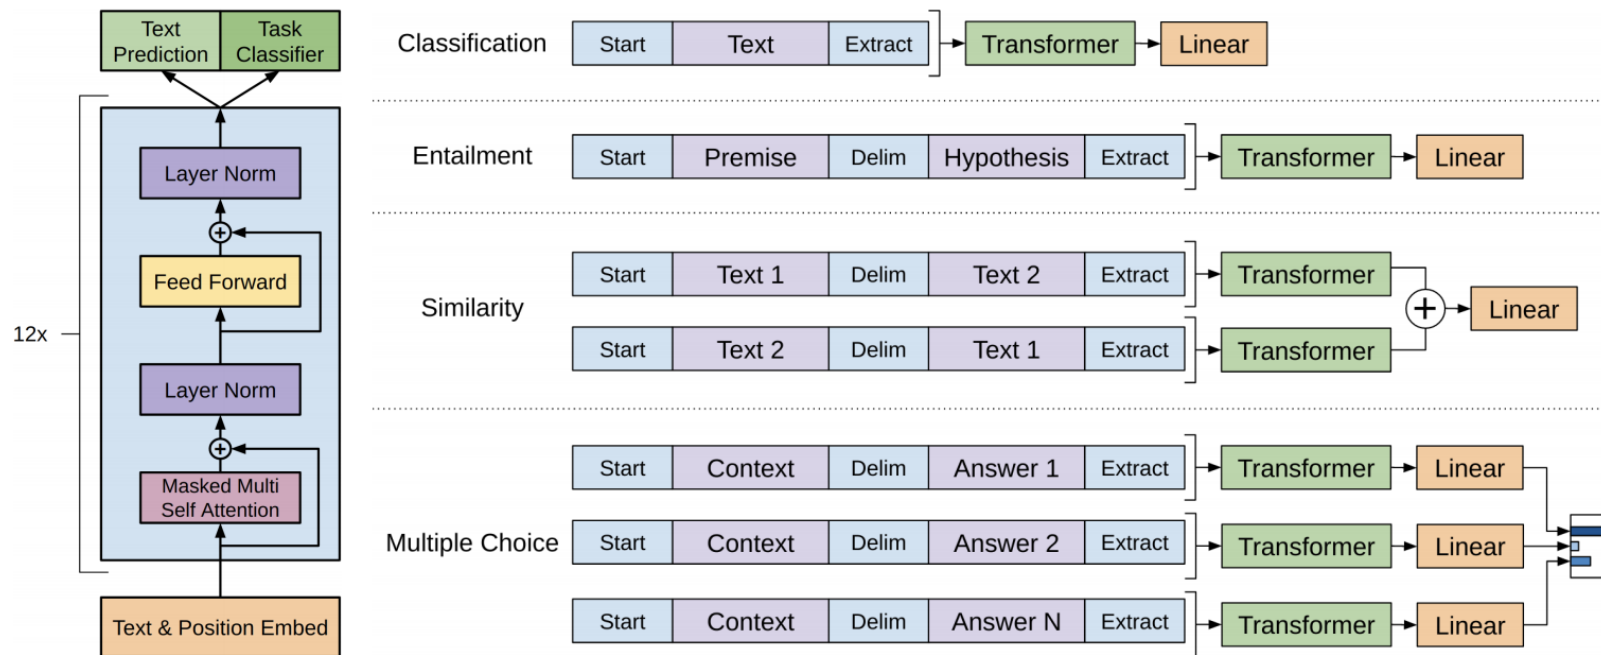
Output Vocabulary: a am I thanks student <eos>



ex. Cross-entropy loss

※ GPT (Generative Pre-Training)

- General model for many NLP tasks
- Only uses Decoder part of the Transformer architecture



References

- Attention Is All You Need (paper): <https://arxiv.org/abs/1706.03762>
- The Illustrated Transformer (blog post): <https://jalammar.github.io/illustrated-transformer/>
- Improving Language Understanding by Generative Pre-Training (paper): https://s3-us-west-2.amazonaws.com/openai-assets/research-covers/language-unsupervised/language_understanding_paper.pdf